

Introduction to Statistical Learning


INF 552, Machine Learning for Data
Informatics

University of Southern California

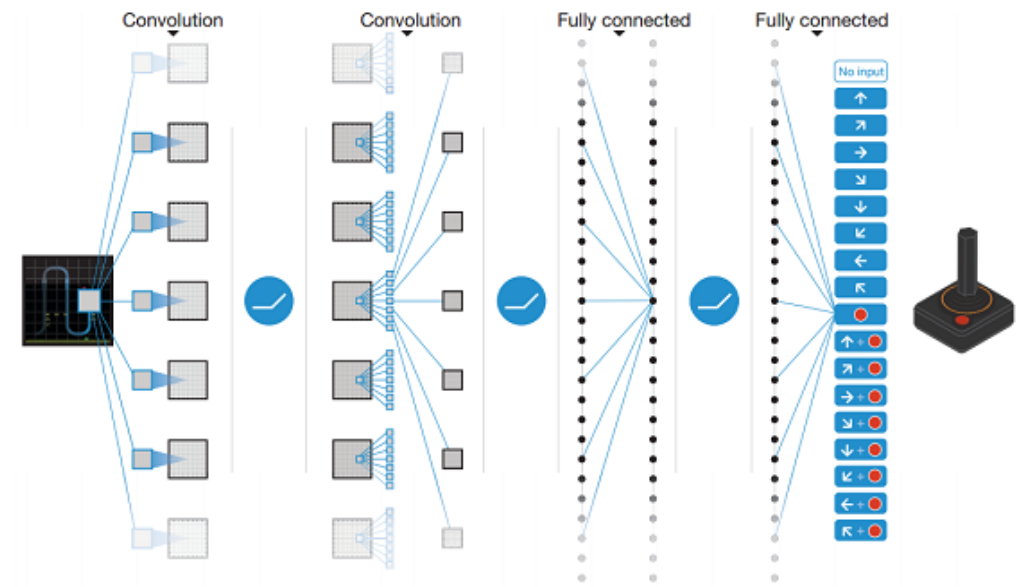
M. R. Rajati, PhD

Lesson 13

Reinforcement Learning



AlphaGo



Overview

- Supervised Learning: Immediate feedback (labels provided for every input).
- Unsupervised Learning: No feedback (no labels provided).
- Reinforcement Learning: Delayed scalar feedback (a number called reward).

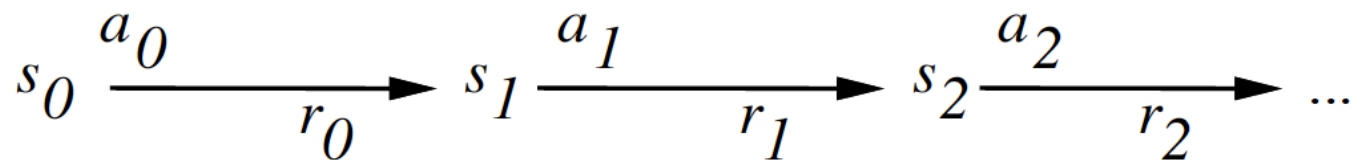
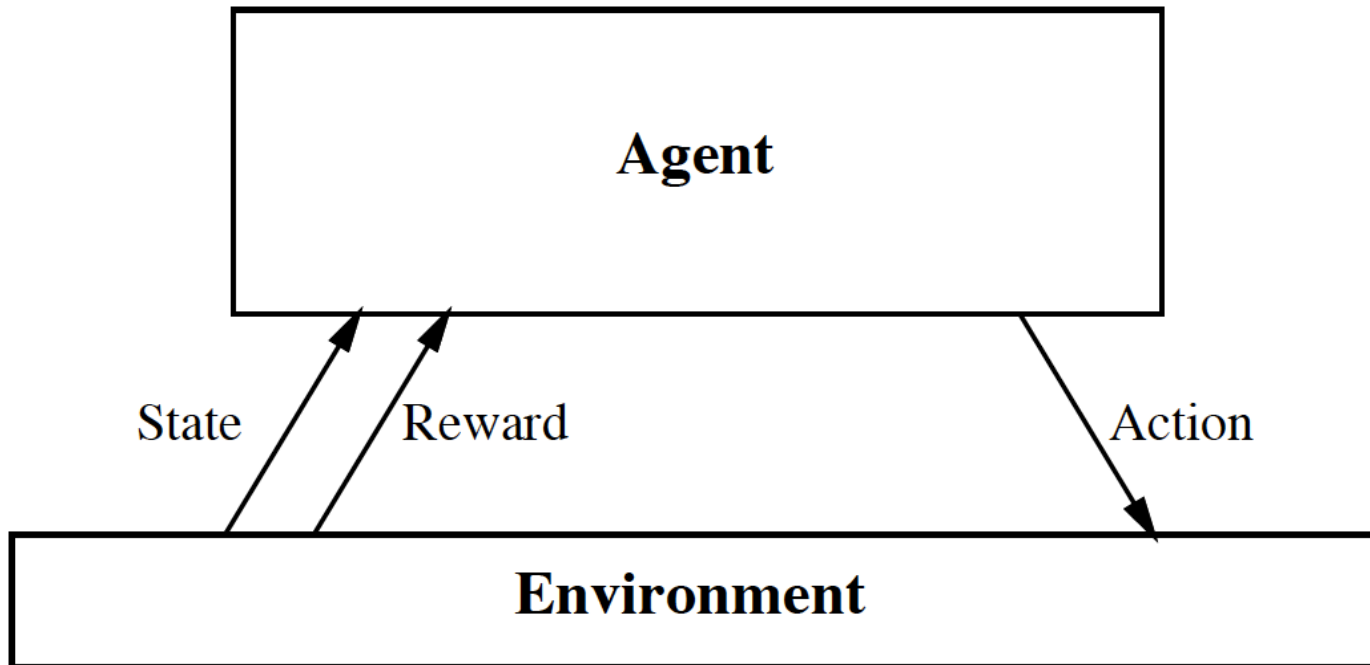
Overview

- RL deals with agents that must sense & act upon their environment.
- This combines classical agent-based AI and machine learning techniques.
It is a very comprehensive problem setting.

Overview

- Examples:
 - A robot cleaning my room and recharging its battery
 - Robot-soccer
 - How to invest in shares
 - Modeling the economy through rational agents
 - Learning how to fly a helicopter
 - Scheduling planes to their destinations
 - and so on

The Big Picture



Your action influences the state of the world which determines its reward

Complications

- The outcome of your actions may be uncertain
- You may not be able to perfectly sense the state of the world
- The reward may be stochastic.
- Reward is delayed (i.e. finding food in a maze)

Complications

- You may have no clue (model) of how rewards are being paid off.
- The world may change while you try to learn it
- How much time do you need to explore uncharted territory before you exploit what you have learned?

The Task

- To learn an optimal *policy* that maps states of the world to actions of the agent.
I.e., if this patch of room is dirty, I clean it. If my battery is empty, I recharge it.

$$\pi : \mathcal{S} \rightarrow \mathcal{A}$$

Action Space
State of Space

- What is it that the agent tries to optimize?
Answer: the **total future discounted reward**:

The Task

du

- What is it that the agent tries to optimize?

Answer: the **total future discounted reward**:

$$V^{\pi}(s_t) = r_t + \gamma r_{t+1} + \gamma^2 r_{t+2} + \dots$$

$$= \sum_{i=0}^{\infty} \gamma^i r_{t+i} \quad 0 \leq \gamma < 1$$

Note: immediate reward is worth more than future reward.

What would happen to mouse in a maze with $\gamma = 0$?

Value Function

- Let's say we have access to the optimal value function that computes the total future discounted reward $V^*(s)$
- What would be the optimal policy $\pi^*(s)$?
- Answer: we choose the action that maximizes:

$$\pi^*(s) = \underset{a}{\operatorname{argmax}} \left[r(s, a) + \underset{t}{\gamma V^*(\delta(s, a))} \right]$$

Value Function

- We assume that we know what the reward will be if we perform action “a” in state “s”:

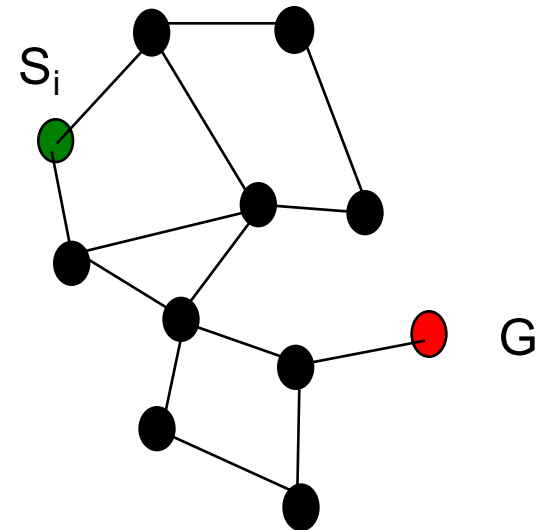
$$r(s, a)$$

- We also assume we know what the next state of the world will be if we perform action “a” in state “s”:

$$s_{t+1} = \delta(s_t, a)$$

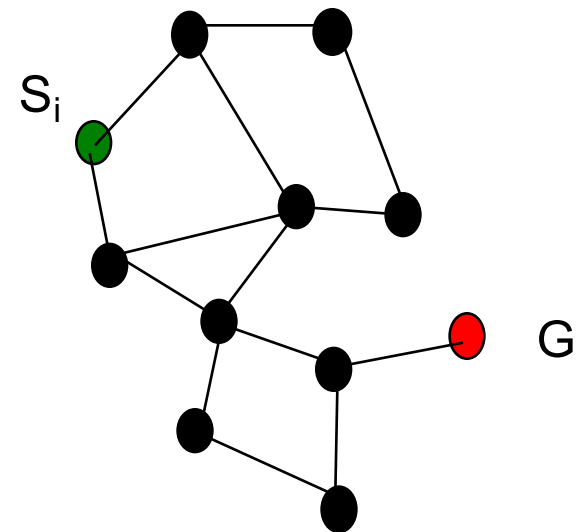
Example I

- Consider some complicated graph, and we would like to find the shortest path from a node S_i to a goal node G .
- Traversing an edge will cost you “length edge” dollars.



Example I

- The value function encodes the total remaining distance to the goal node from any node s , i.e.
 $V(s) = \text{“}1 / \textit{distance}\text{”}$ to goal from s .
- If you know $V(s)$, the problem is trivial. You simply choose the node that has highest $V(s)$.




Example II


- A simple deterministic world.
- Each grid square represents a distinct state, each arrow a distinct action.
- The immediate reward function, $r(s, a)$ gives reward 100 for actions entering the goal state G , and zero otherwise. Values of $V^*(s)$ follow from $r(s, a)$, and the discount factor $\gamma = 0.9$.
- An optimal policy is also shown.

Example II


Find your way to the goal.



$r(s, a)$ (immediate reward) values



$V^*(s)$ values



One optimal policy

$$V^\pi(s_t) = r_t + \gamma r_{t+1} + \gamma^2 r_{t+2} + \dots$$

$$= \sum_{i=0}^{\infty} \gamma^i r_{t+i} \quad 0 \leq \gamma < 1$$

Q-Function

- One approach to RL is then to try to estimate $V^*(s)$.

Bellman Equation:

$$V^*(s) \leftarrow \max_a \left[r(s,a) + \gamma V^*(\delta(s,a)) \right]$$

- However, this approach requires you to know $r(s,a)$ and $\delta(s,a)$.
- This is unrealistic in many real problems. What is the reward if a robot is exploring mars and decides to take a right turn?

Q-Function

- Fortunately we can circumvent this problem by exploring and experiencing how the world reacts to our actions. We need to *learn* r & δ .

Q-Function

- We want a function that directly learns good state-action pairs, i.e. what action should I take in this state. We call this $Q(s,a)$.

Q-Function

- Let us define the evaluation function $Q(s, a)$ so that its value is the maximum discounted cumulative reward that can be achieved starting from state s and applying action a as the first action.
- In other words, the value of Q is the reward received immediately upon executing action a from state s , plus the value (discounted by γ) of following the optimal policy thereafter.

$$Q(s, a) \equiv r(s, a) + \gamma V^*(\delta(s, a))$$

Q-Function

- Why is this rewrite important? Because it shows that if the agent learns the Q function instead of the V^* function, it will be able to select optimal actions even when it has no knowledge of the functions r and δ .

Q-Function

- It need only consider each available action a in its current state s and choose the action that maximizes $Q(s, a)$.

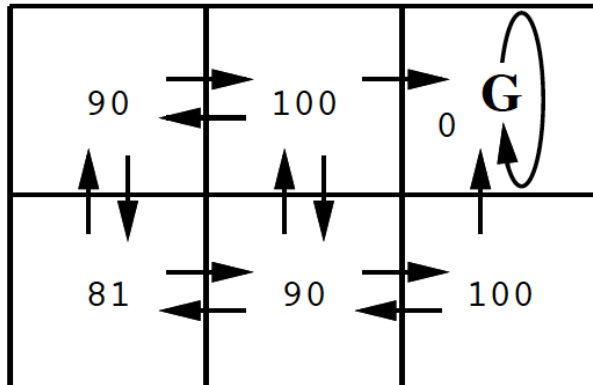
$$\pi^*(s) = \underset{a}{\operatorname{argmax}} Q(s, a)$$

$$V^*(s) = \max_a Q(s, a)$$

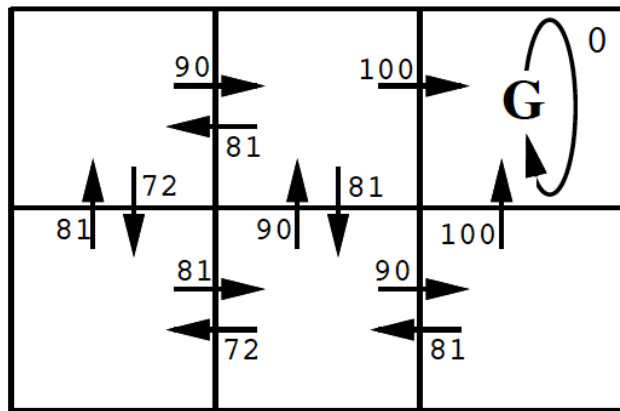
Example

- To illustrate, the figure in the next slide shows the Q values for every state and action in the simple grid world.
- The Q value for each state-action transition equals the r value for this transition plus the V^* value for the resulting state discounted by γ .
- The optimal policy shown in the figure corresponds to selecting actions with maximal Q values.

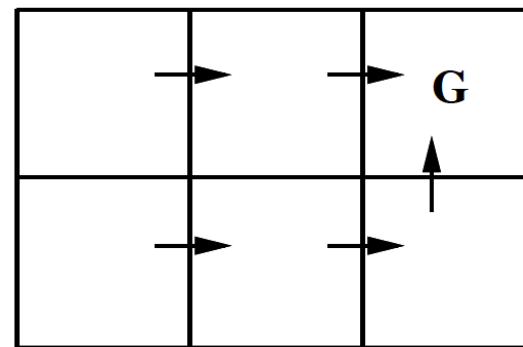
Example II



$V^*(s)$ values



$Q(s, a)$ values



One optimal policy

Check that

$$\pi^*(s) = \underset{a}{\operatorname{argmax}} Q(s, a)$$

$$V^*(s) = \max_a Q(s, a)$$

Q-Learning

- Learning the Q function corresponds to learning the optimal policy. How can Q be learned?
- The key problem is finding a reliable way to estimate training values for Q , given only a sequence of immediate rewards r spread out over time.

Q-Learning

- This can be accomplished through iterative approximation.
- To see how, notice the close relationship between Q and V^* ,

$$V^*(s) = \max_{a'} Q(s, a')$$

Q-Learning

which allows rewriting the Q function as:

$$\begin{aligned} Q(s, a) &\equiv r(s, a) + \gamma V^*(\delta(s, a)) \\ &= r(s, a) + \gamma \max_{a'} Q(\delta(s, a), a') \end{aligned}$$

This still depends on $r(s, a)$ and $\delta(s, a)$;
however,

Q-Learning

$$\begin{aligned} Q(s,a) &\equiv r(s,a) + \gamma V^*(\delta(s,a)) \\ &= r(s,a) + \gamma \max_{a'} Q(\delta(s,a), a') \end{aligned}$$

this recursive definition of Q provides the basis for algorithms that iteratively approximate Q .

\hat{Q} refers to the learner's estimate, or hypothesis, of the actual Q function.

Q-Learning

$$\begin{aligned} Q(s, a) &\equiv r(s, a) + \gamma V^*(\delta(s, a)) \\ &= r(s, a) + \gamma \max_{a'} Q(\delta(s, a), a') \end{aligned}$$

- Imagine the robot is exploring its environment, trying new actions as it goes.
- At every step it receives some reward “ r ”, and it observes the environment change into a new state s' for action a .
- How can we use these observations, (s, a, s', r) to learn a model?

Q-Learning

- The learner represents its hypothesis \hat{Q} by a large table with a separate entry for each state-action pair.
- The table entry for the pair (s, a) stores the value for $\hat{Q}(s,a)$, learner's current hypothesis about the actual but unknown value $Q(s,a)$.
- The table can be initially filled with random values (though it is easier to understand the algorithm if one assumes initial values of zero).

Q-Learning

- The agent repeatedly observes its current state s , chooses some action a , executes this action, then observes the resulting reward $r = r(s, a)$ and the new state $s' = \delta(s, a)$.

Q-Learning

- It then updates the table entry for $\hat{Q}(s,a)$ following each such transition, according to the rule:

$$\hat{Q}(s,a) \leftarrow r + \gamma \max_{a'} \hat{Q}(s',a') \quad s' = s_{t+1}$$

Q-Learning

$$\hat{Q}(s, a) \leftarrow r + \gamma \max_{a'} \hat{Q}(s', a') \quad s' = s_{t+1}$$

- This equation continually estimates Q at state s consistent with an estimate of Q at state s' , one step in the future: temporal difference (TD) learning.
- Note that s' is closer to goal, and hence more “reliable”, but still an estimate itself.

Q-Learning Summary

Q learning algorithm

For each s, a initialize the table entry $\hat{Q}(s, a)$ to zero.

Observe the current state s

Do forever:

- Select an action a and execute it
- Receive immediate reward r
- Observe the new state s'
- Update the table entry for $\hat{Q}(s, a)$ as follows:

$$\hat{Q}(s, a) \leftarrow r + \gamma \max_{a'} \hat{Q}(s', a')$$

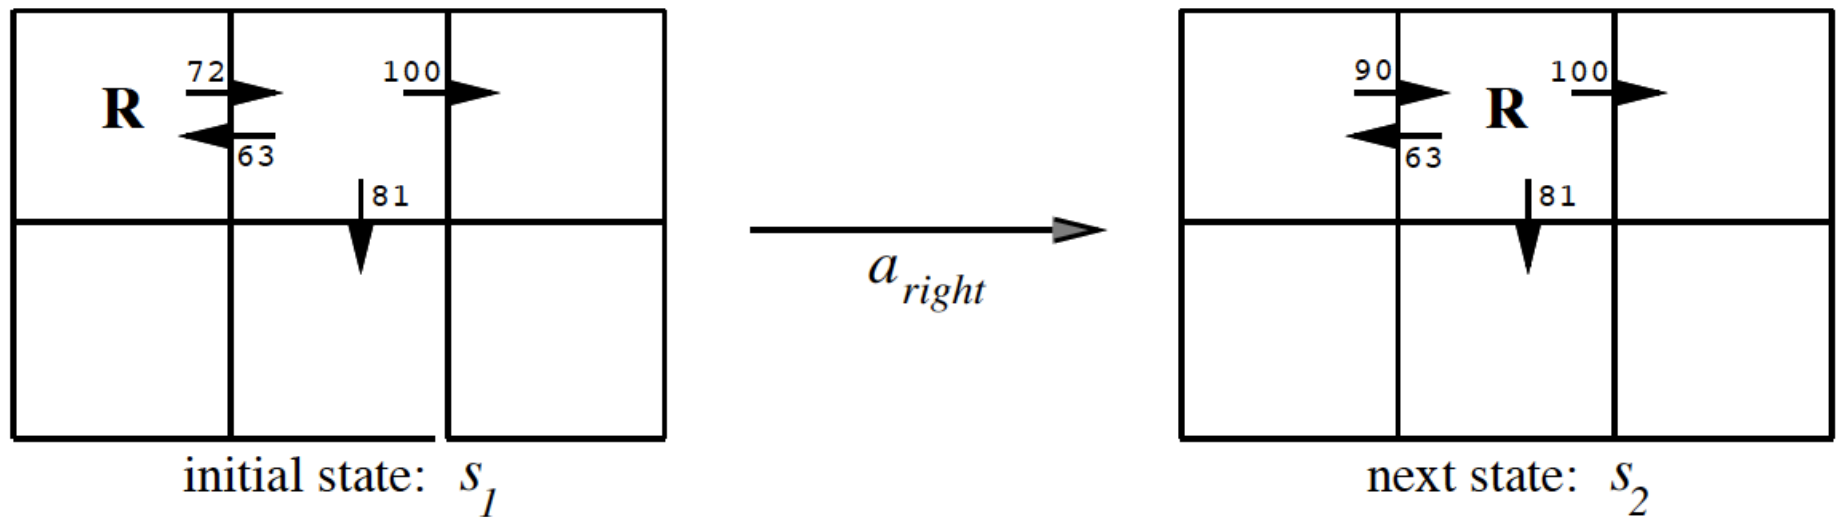
- $s \leftarrow s'$

Q-Learning

$$\hat{Q}(s, a) \leftarrow r + \gamma \max_{a'} \hat{Q}(s', a')$$
$$s' = s_{t+1}$$

- We do an update after each state-action pair. I.e., we are learning online!
- We are learning useful things about explored state-action pairs. These are typically most useful because they are likely to be encountered again.
- Under suitable conditions, these updates can actually be proved to converge to the real answer.

Example: Q-Learning



$$\begin{aligned}\hat{Q}(s_1, a_{right}) &\leftarrow r + \gamma \max_{a'} \hat{Q}(s_2, a') \\ &\leftarrow 0 + 0.9 \max\{66, 81, 100\} \\ &\leftarrow 90\end{aligned}$$

Q-learning propagates Q-estimates 1-step backwards